

```
In [1]: #Import required Libraries
import torch
import numpy as np
import torch.nn as nn
import pandas as pd
import tensorflow as tf
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: #Create dataframe to store our dataset
loan_df = pd.read_csv("Dropbox/Coding/AI and ML Bootcamp/Course 4 Deep Learning with K
```

```
In [3]: #Feature Transformation
loan_df = pd.get_dummies(loan_df, columns=['purpose'], drop_first=True)
print(loan_df.head())
```

	credit.policy	int.rate	installment	log.annual.inc	dti	fico	\
0	1	0.1189	829.10	11.350407	19.48	737	
1	1	0.1071	228.22	11.082143	14.29	707	
2	1	0.1357	366.86	10.373491	11.63	682	
3	1	0.1008	162.34	11.350407	8.10	712	
4	1	0.1426	102.92	11.299732	14.97	667	

	days.with.cr.line	revol.bal	revol.util	inq.last.6mths	delinq.2yrs	\
0	5639.958333	28854	52.1	0	0	
1	2760.000000	33623	76.7	0	0	
2	4710.000000	3511	25.6	1	0	
3	2699.958333	33667	73.2	1	0	
4	4066.000000	4740	39.5	0	1	

	pub.rec	not.fully.paid	purpose_credit_card	purpose_debt_consolidation	\
0	0	0	0	1	
1	0	0	1	0	
2	0	0	0	1	
3	0	0	0	1	
4	0	0	1	0	

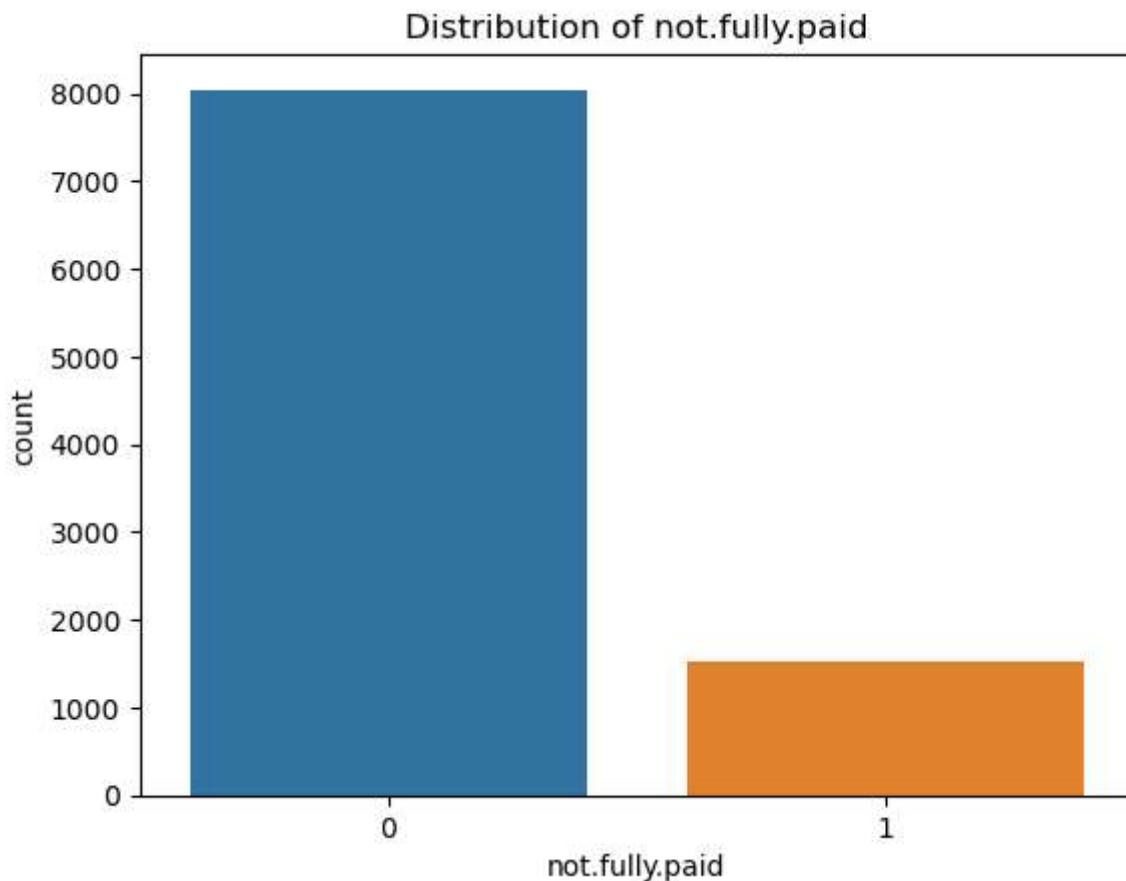
	purpose_educational	purpose_home_improvement	purpose_major_purchase	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	

	purpose_small_business
0	0
1	0
2	0
3	0
4	0

```
In [4]: #Exploratory data analysis
#Look at descriptions of variables
print(loan_df.describe())
```

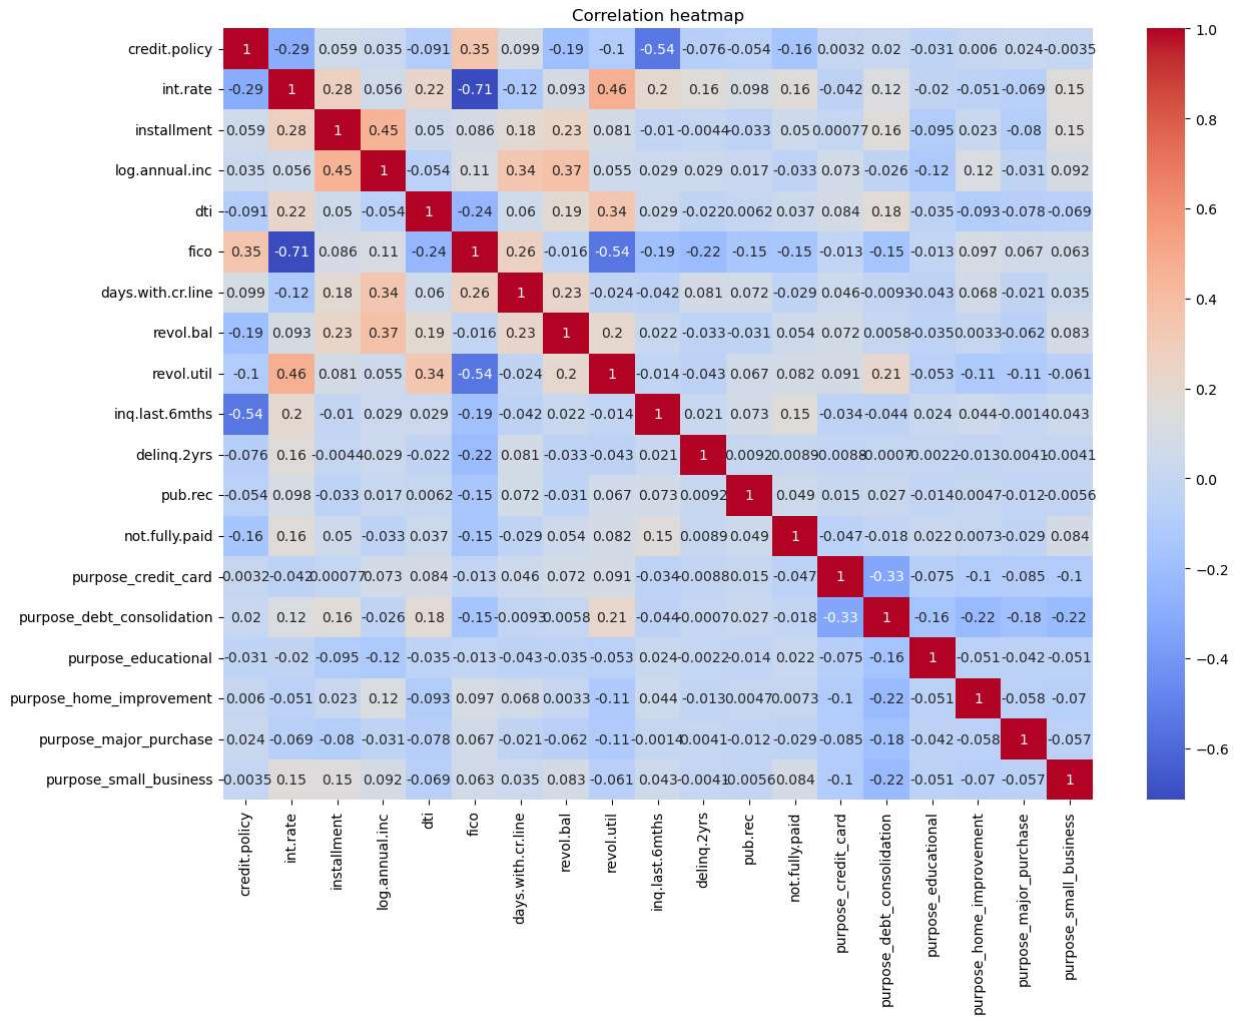
	credit.policy	int.rate	installment	log.annual.inc	dti	\
count	9578.000000	9578.000000	9578.000000	9578.000000	9578.000000	
mean	0.804970	0.122640	319.089413	10.932117	12.606679	
std	0.396245	0.026847	207.071301	0.614813	6.883970	
min	0.000000	0.060000	15.670000	7.547502	0.000000	
25%	1.000000	0.103900	163.770000	10.558414	7.212500	
50%	1.000000	0.122100	268.950000	10.928884	12.665000	
75%	1.000000	0.140700	432.762500	11.291293	17.950000	
max	1.000000	0.216400	940.140000	14.528354	29.960000	
	fico	days.with.cr.line	revol.bal	revol.util	\	
count	9578.000000	9578.000000	9.578000e+03	9578.000000		
mean	710.846314	4560.767197	1.691396e+04	46.799236		
std	37.970537	2496.930377	3.375619e+04	29.014417		
min	612.000000	178.958333	0.000000e+00	0.000000		
25%	682.000000	2820.000000	3.187000e+03	22.600000		
50%	707.000000	4139.958333	8.596000e+03	46.300000		
75%	737.000000	5730.000000	1.824950e+04	70.900000		
max	827.000000	17639.958330	1.207359e+06	119.000000		
	inq.last.6mths	delinq.2yrs	pub.rec	not.fully.paid	\	
count	9578.000000	9578.000000	9578.000000	9578.000000		
mean	1.577469	0.163708	0.062122	0.160054		
std	2.200245	0.546215	0.262126	0.366676		
min	0.000000	0.000000	0.000000	0.000000		
25%	0.000000	0.000000	0.000000	0.000000		
50%	1.000000	0.000000	0.000000	0.000000		
75%	2.000000	0.000000	0.000000	0.000000		
max	33.000000	13.000000	5.000000	1.000000		
	purpose_credit_card	purpose_debt_consolidation	purpose_educational	\		
count	9578.000000	9578.000000	9578.000000	9578.000000		
mean	0.131760	0.413134	0.035811			
std	0.338248	0.492422	0.185829			
min	0.000000	0.000000	0.000000			
25%	0.000000	0.000000	0.000000			
50%	0.000000	0.000000	0.000000			
75%	0.000000	1.000000	0.000000			
max	1.000000	1.000000	1.000000			
	purpose_home_improvement	purpose_major_purchase	\			
count	9578.000000	9578.000000				
mean	0.065671	0.045625				
std	0.247720	0.208682				
min	0.000000	0.000000				
25%	0.000000	0.000000				
50%	0.000000	0.000000				
75%	0.000000	0.000000				
max	1.000000	1.000000				
	purpose_small_business					
count	9578.000000					
mean	0.064627					
std	0.245880					
min	0.000000					
25%	0.000000					
50%	0.000000					
75%	0.000000					
max	1.000000					

```
In [5]: #Check distribution of target variable  
sns.countplot(x='not.fully.paid', data=loan_df)  
plt.title('Distribution of not.fully.paid')  
plt.show()
```



```
In [6]: #Create correlation matrix to check variable correlations  
correlation_matrix = loan_df.corr()  
plt.figure(figsize=(14,10))  
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')  
plt.title('Correlation heatmap')  
plt.show()
```

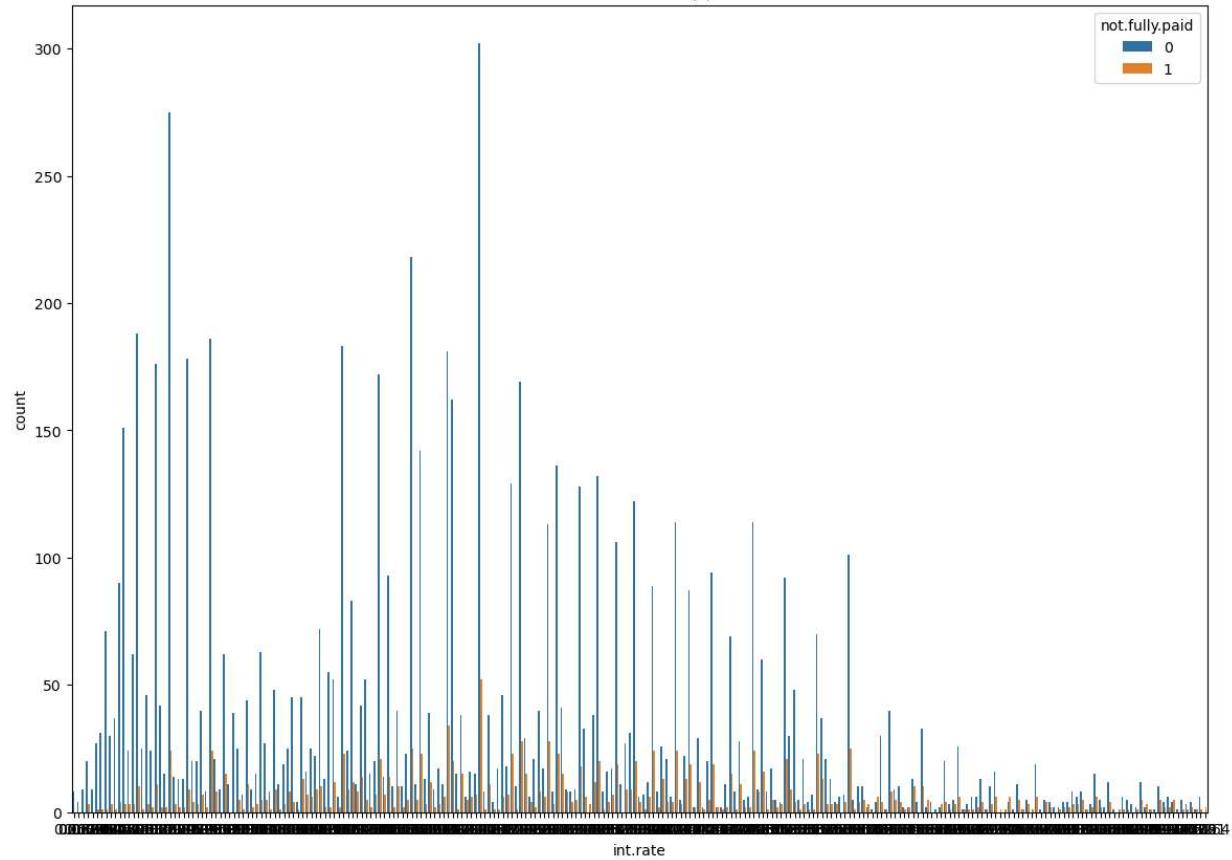
Lending Club Loan Project



```
In [7]: #int.rate and inq.last.6mths are most closely correlated to not.fully.paid
#Features with highest correlation are int.rate and revol.util, installment and Log.ar
```

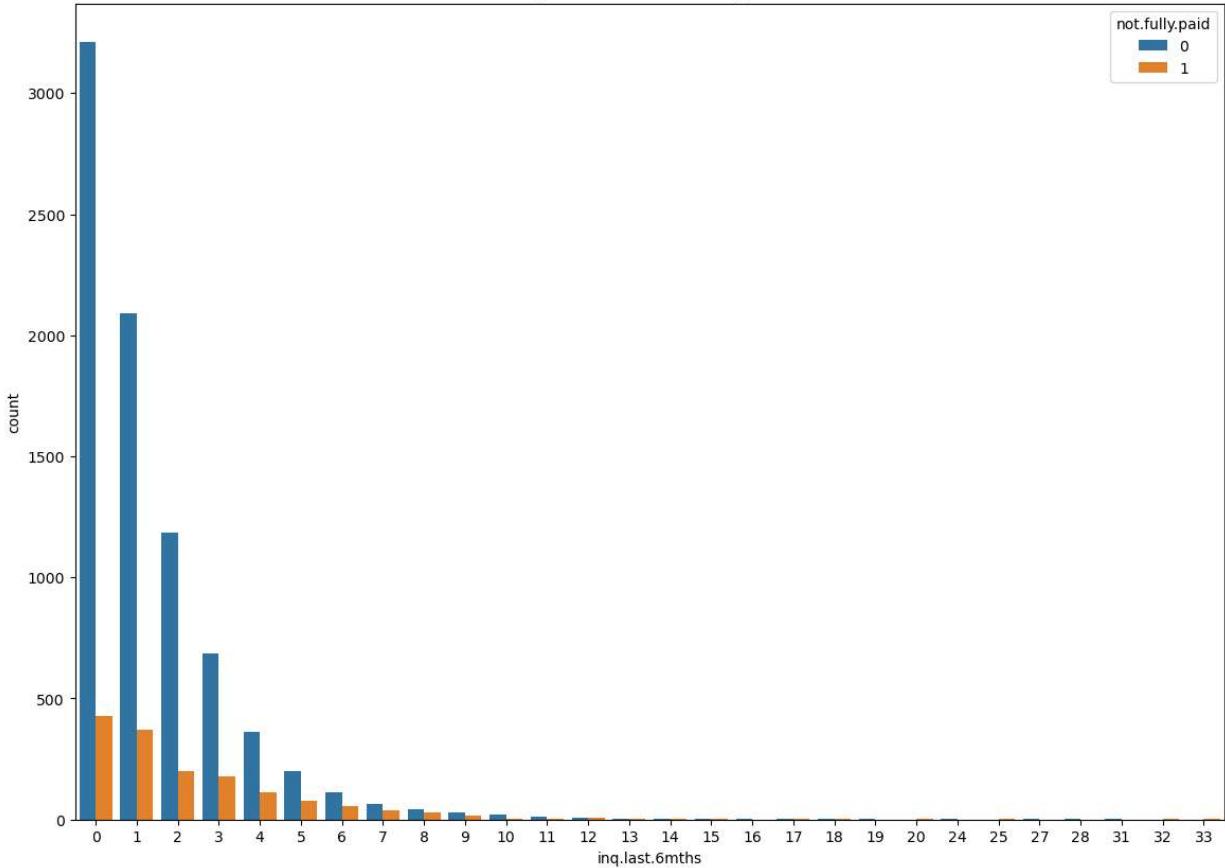
```
In [8]: #'int.rate' variable correlation with not.fully.paid
plt.figure(figsize=(14,10))
sns.countplot(x='int.rate', hue='not.fully.paid', data=loan_df)
plt.title('int.rate vs not.fully.paid')
plt.show()
```

int.rate vs not.fully.paid



```
In [9]: #'inq.Last.6mths' variable correlation with not.fully.paid  
plt.figure(figsize=(14,10))  
sns.countplot(x='inq.last.6mths', hue='not.fully.paid', data=loan_df)  
plt.title('inq.last.6mths vs not.fully.paid')  
plt.show()
```

inq.last.6mths vs not.fully.paid



```
In [10]: #Additional feature engineering
#Drop features with low correlation to target variable. The least correlated features
loan_df_mod = loan_df.drop('credit.policy', axis=1)
loan_df_mod = loan_df_mod.drop('fico', axis=1)
```

```
In [11]: #Drop 'log.annual.inc' due to low correlation with target and high correlation with 'i
loan_df_mod = loan_df_mod.drop('log.annual.inc', axis=1)
```

```
In [12]: #Prepare data for input into model
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
X = loan_df_mod.drop('not.fully.paid', axis=1)
y = loan_df_mod['not.fully.paid']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
scaler = StandardScaler().fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [72]: #Build the model
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.regularizers import l2
model = Sequential()
#Input Layer
model.add(Dense(128, activation='relu', input_dim=X_train.shape[1]))
model.add(Dropout(0.3))
#Hidden Layer 1 (with added weight regularization to help fix overfitting)
model.add(Dense(64, activation='relu', kernel_regularizer=l2(0.001)))
model.add(Dropout(0.3))
#Hidden Layer 2
```

```
model.add(Dense(64, activation='relu'))
model.add(Dropout(0.3))
#Output Layer
model.add(Dense(1, activation='sigmoid'))
#Add learning rate to optimizer to reduce overfitting
optimizer = tf.keras.optimizers.Adam(learning_rate=0.01)
#Compile model
model.compile(optimizer=optimizer, loss='binary_crossentropy', metrics=['accuracy'])
```

In [73]: #View model summary
model.summary()

Model: "sequential_16"

Layer (type)	Output Shape	Param #
<hr/>		
dense_57 (Dense)	(None, 128)	2048
dropout_41 (Dropout)	(None, 128)	0
dense_58 (Dense)	(None, 64)	8256
dropout_42 (Dropout)	(None, 64)	0
dense_59 (Dense)	(None, 64)	4160
dropout_43 (Dropout)	(None, 64)	0
dense_60 (Dense)	(None, 1)	65
<hr/>		
Total params: 14529 (56.75 KB)		
Trainable params: 14529 (56.75 KB)		
Non-trainable params: 0 (0.00 Byte)		

In [74]: #Train the model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=400, ba

```
Epoch 1/400
240/240 [=====] - 1s 1ms/step - loss: 0.4865 - accuracy: 0.8
370 - val_loss: 0.4297 - val_accuracy: 0.8408
Epoch 2/400
240/240 [=====] - 0s 933us/step - loss: 0.4401 - accuracy:
0.8395 - val_loss: 0.4188 - val_accuracy: 0.8408
Epoch 3/400
240/240 [=====] - 0s 937us/step - loss: 0.4353 - accuracy:
0.8393 - val_loss: 0.4271 - val_accuracy: 0.8408
Epoch 4/400
240/240 [=====] - 0s 921us/step - loss: 0.4337 - accuracy:
0.8395 - val_loss: 0.4258 - val_accuracy: 0.8408
Epoch 5/400
240/240 [=====] - 0s 929us/step - loss: 0.4342 - accuracy:
0.8393 - val_loss: 0.4219 - val_accuracy: 0.8408
Epoch 6/400
240/240 [=====] - 0s 920us/step - loss: 0.4339 - accuracy:
0.8389 - val_loss: 0.4190 - val_accuracy: 0.8408
Epoch 7/400
240/240 [=====] - 0s 916us/step - loss: 0.4302 - accuracy:
0.8393 - val_loss: 0.4179 - val_accuracy: 0.8408
Epoch 8/400
240/240 [=====] - 0s 912us/step - loss: 0.4326 - accuracy:
0.8396 - val_loss: 0.4190 - val_accuracy: 0.8408
Epoch 9/400
240/240 [=====] - 0s 912us/step - loss: 0.4285 - accuracy:
0.8391 - val_loss: 0.4176 - val_accuracy: 0.8408
Epoch 10/400
240/240 [=====] - 0s 912us/step - loss: 0.4323 - accuracy:
0.8399 - val_loss: 0.4255 - val_accuracy: 0.8408
Epoch 11/400
240/240 [=====] - 0s 912us/step - loss: 0.4310 - accuracy:
0.8397 - val_loss: 0.4216 - val_accuracy: 0.8408
Epoch 12/400
240/240 [=====] - 0s 922us/step - loss: 0.4334 - accuracy:
0.8397 - val_loss: 0.4229 - val_accuracy: 0.8408
Epoch 13/400
240/240 [=====] - 0s 921us/step - loss: 0.4310 - accuracy:
0.8396 - val_loss: 0.4228 - val_accuracy: 0.8408
Epoch 14/400
240/240 [=====] - 0s 933us/step - loss: 0.4286 - accuracy:
0.8392 - val_loss: 0.4193 - val_accuracy: 0.8408
Epoch 15/400
240/240 [=====] - 0s 941us/step - loss: 0.4306 - accuracy:
0.8396 - val_loss: 0.4203 - val_accuracy: 0.8408
Epoch 16/400
240/240 [=====] - 0s 954us/step - loss: 0.4293 - accuracy:
0.8396 - val_loss: 0.4212 - val_accuracy: 0.8408
Epoch 17/400
240/240 [=====] - 0s 929us/step - loss: 0.4303 - accuracy:
0.8395 - val_loss: 0.4363 - val_accuracy: 0.8408
Epoch 18/400
240/240 [=====] - 0s 921us/step - loss: 0.4341 - accuracy:
0.8397 - val_loss: 0.4267 - val_accuracy: 0.8408
Epoch 19/400
240/240 [=====] - 0s 921us/step - loss: 0.4296 - accuracy:
0.8397 - val_loss: 0.4186 - val_accuracy: 0.8408
Epoch 20/400
240/240 [=====] - 0s 929us/step - loss: 0.4317 - accuracy:
0.8397 - val_loss: 0.4173 - val_accuracy: 0.8408
```

```
Epoch 361/400
240/240 [=====] - 0s 916us/step - loss: 0.4336 - accuracy: 0.8397 - val_loss: 0.4278 - val_accuracy: 0.8408
Epoch 362/400
240/240 [=====] - 0s 921us/step - loss: 0.4333 - accuracy: 0.8397 - val_loss: 0.4253 - val_accuracy: 0.8408
Epoch 363/400
240/240 [=====] - 0s 920us/step - loss: 0.4372 - accuracy: 0.8397 - val_loss: 0.4313 - val_accuracy: 0.8408
Epoch 364/400
240/240 [=====] - 0s 916us/step - loss: 0.4385 - accuracy: 0.8397 - val_loss: 0.4350 - val_accuracy: 0.8408
Epoch 365/400
240/240 [=====] - 0s 967us/step - loss: 0.4378 - accuracy: 0.8397 - val_loss: 0.4274 - val_accuracy: 0.8408
Epoch 366/400
240/240 [=====] - 0s 971us/step - loss: 0.4401 - accuracy: 0.8397 - val_loss: 0.4299 - val_accuracy: 0.8408
Epoch 367/400
240/240 [=====] - 0s 983us/step - loss: 0.4343 - accuracy: 0.8397 - val_loss: 0.4293 - val_accuracy: 0.8408
Epoch 368/400
240/240 [=====] - 0s 992us/step - loss: 0.4350 - accuracy: 0.8397 - val_loss: 0.4266 - val_accuracy: 0.8408
Epoch 369/400
240/240 [=====] - 0s 975us/step - loss: 0.4347 - accuracy: 0.8397 - val_loss: 0.4231 - val_accuracy: 0.8408
Epoch 370/400
240/240 [=====] - 0s 975us/step - loss: 0.4326 - accuracy: 0.8397 - val_loss: 0.4291 - val_accuracy: 0.8408
Epoch 371/400
240/240 [=====] - 0s 933us/step - loss: 0.4368 - accuracy: 0.8397 - val_loss: 0.4240 - val_accuracy: 0.8408
Epoch 372/400
240/240 [=====] - 0s 923us/step - loss: 0.4347 - accuracy: 0.8397 - val_loss: 0.4270 - val_accuracy: 0.8408
Epoch 373/400
240/240 [=====] - 0s 962us/step - loss: 0.4311 - accuracy: 0.8397 - val_loss: 0.4302 - val_accuracy: 0.8408
Epoch 374/400
240/240 [=====] - 0s 962us/step - loss: 0.4324 - accuracy: 0.8397 - val_loss: 0.4302 - val_accuracy: 0.8408
Epoch 375/400
240/240 [=====] - 0s 967us/step - loss: 0.4325 - accuracy: 0.8397 - val_loss: 0.4297 - val_accuracy: 0.8408
Epoch 376/400
240/240 [=====] - 0s 925us/step - loss: 0.4328 - accuracy: 0.8397 - val_loss: 0.4272 - val_accuracy: 0.8408
Epoch 377/400
240/240 [=====] - 0s 1ms/step - loss: 0.4324 - accuracy: 0.8397 - val_loss: 0.4360 - val_accuracy: 0.8408
Epoch 378/400
240/240 [=====] - 0s 971us/step - loss: 0.4360 - accuracy: 0.8397 - val_loss: 0.4282 - val_accuracy: 0.8408
Epoch 379/400
240/240 [=====] - 0s 933us/step - loss: 0.4364 - accuracy: 0.8397 - val_loss: 0.4267 - val_accuracy: 0.8408
Epoch 380/400
240/240 [=====] - 0s 946us/step - loss: 0.4306 - accuracy: 0.8397 - val_loss: 0.4235 - val_accuracy: 0.8408
```

```
Epoch 381/400
240/240 [=====] - 0s 912us/step - loss: 0.4350 - accuracy: 0.8397 - val_loss: 0.4333 - val_accuracy: 0.8408
Epoch 382/400
240/240 [=====] - 0s 916us/step - loss: 0.4371 - accuracy: 0.8397 - val_loss: 0.4339 - val_accuracy: 0.8408
Epoch 383/400
240/240 [=====] - 0s 971us/step - loss: 0.4377 - accuracy: 0.8397 - val_loss: 0.4366 - val_accuracy: 0.8408
Epoch 384/400
240/240 [=====] - 0s 920us/step - loss: 0.4367 - accuracy: 0.8397 - val_loss: 0.4246 - val_accuracy: 0.8408
Epoch 385/400
240/240 [=====] - 0s 921us/step - loss: 0.4378 - accuracy: 0.8397 - val_loss: 0.4271 - val_accuracy: 0.8408
Epoch 386/400
240/240 [=====] - 0s 962us/step - loss: 0.4340 - accuracy: 0.8397 - val_loss: 0.4254 - val_accuracy: 0.8408
Epoch 387/400
240/240 [=====] - 0s 925us/step - loss: 0.4349 - accuracy: 0.8397 - val_loss: 0.4305 - val_accuracy: 0.8408
Epoch 388/400
240/240 [=====] - 0s 916us/step - loss: 0.4382 - accuracy: 0.8397 - val_loss: 0.4272 - val_accuracy: 0.8408
Epoch 389/400
240/240 [=====] - 0s 926us/step - loss: 0.4337 - accuracy: 0.8397 - val_loss: 0.4295 - val_accuracy: 0.8408
Epoch 390/400
240/240 [=====] - 0s 912us/step - loss: 0.4340 - accuracy: 0.8397 - val_loss: 0.4277 - val_accuracy: 0.8408
Epoch 391/400
240/240 [=====] - 0s 916us/step - loss: 0.4347 - accuracy: 0.8397 - val_loss: 0.4294 - val_accuracy: 0.8408
Epoch 392/400
240/240 [=====] - 0s 912us/step - loss: 0.4325 - accuracy: 0.8397 - val_loss: 0.4329 - val_accuracy: 0.8408
Epoch 393/400
240/240 [=====] - 0s 917us/step - loss: 0.4399 - accuracy: 0.8397 - val_loss: 0.4333 - val_accuracy: 0.8408
Epoch 394/400
240/240 [=====] - 0s 908us/step - loss: 0.4373 - accuracy: 0.8397 - val_loss: 0.4342 - val_accuracy: 0.8408
Epoch 395/400
240/240 [=====] - 0s 975us/step - loss: 0.4348 - accuracy: 0.8397 - val_loss: 0.4304 - val_accuracy: 0.8408
Epoch 396/400
240/240 [=====] - 0s 912us/step - loss: 0.4383 - accuracy: 0.8397 - val_loss: 0.4320 - val_accuracy: 0.8408
Epoch 397/400
240/240 [=====] - 0s 921us/step - loss: 0.4391 - accuracy: 0.8397 - val_loss: 0.4372 - val_accuracy: 0.8408
Epoch 398/400
240/240 [=====] - 0s 916us/step - loss: 0.4348 - accuracy: 0.8397 - val_loss: 0.4312 - val_accuracy: 0.8408
Epoch 399/400
240/240 [=====] - 0s 921us/step - loss: 0.4343 - accuracy: 0.8397 - val_loss: 0.4281 - val_accuracy: 0.8408
Epoch 400/400
240/240 [=====] - 0s 962us/step - loss: 0.4341 - accuracy: 0.8397 - val_loss: 0.4290 - val_accuracy: 0.8408
```

In [75]: #Evaluate the matrix

```
loss, accuracy = model.evaluate(X_test, y_test)
print(f"Test Loss: {loss:.4f}")
print(f"Test Accuracy: {accuracy:.4f}")
```

60/60 [=====] - 0s 475us/step - loss: 0.4290 - accuracy: 0.8

408

Test Loss: 0.4290

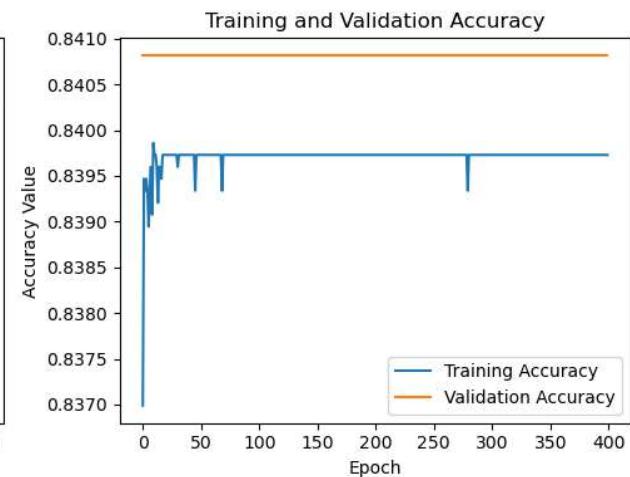
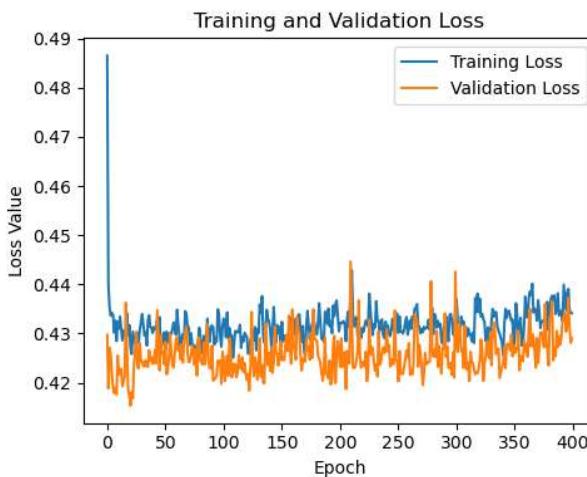
Test Accuracy: 0.8408

In [78]: #Plot training and validation Loss

```
plt.figure(figsize=(10,4))
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss Value')
plt.legend()

#Plot training and validation accuracy
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy Value')
plt.legend()

plt.tight_layout()
plt.show()
```



During exploratory data analysis and feature engineering, the variables 'credit.policy' and 'fico' were dropped because they had the lowest correlation with the target variable 'not.fully.paid'. The variable 'log.annual.inc' was dropped because it had a high correlation with 'installment' and a relatively low correlation with the target. The model was initially built with one hidden layer, an unspecified learning rate and no regularizer attached. The results heavily suggested overfitting, and the model had a low initial validation accuracy and high validation loss. A regularizer was added to the hidden layer to reduce overfitting and multiple values were tested to find the best result. The optimizer was then given a learning rate, and 0.1 was found to be

the best value. After the validation results were better but still somewhat unsatisfying, another two hidden layers were added, but only two hidden layers were kept because having three yielded unfavorable results.

In []: